



```

import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import precision_recall_curve
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.preprocessing import MinMaxScaler

from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

!ls

drive sample_data

df=pd.read_csv("/content/drive/MyDrive/logistic_regression.csv")

df.head()

```

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years

```

5 rows × 9 columns

df.shape

(396030, 27)

df["loan_status"].value_counts(normalize=True) * 100

Fully Paid      80.387092
Charged Off     19.612908
Name: loan_status, dtype: float64

df.describe(include='all')

```

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title
count	396030.000000	396030	396030.000000	396030.000000	396030	396030	373
unique	NaN	2	NaN	NaN	7	35	173
top	NaN	36 months	NaN	NaN	B	B3	Teacher
freq	NaN	302005	NaN	NaN	116018	26655	4
mean	14113.888089	NaN	13.639400	431.849698	NaN	NaN	1
std	8357.441341	NaN	4.472157	250.727790	NaN	NaN	1
min	500.000000	NaN	5.320000	16.080000	NaN	NaN	1
25%	8000.000000	NaN	10.490000	250.330000	NaN	NaN	1
50%	12000.000000	NaN	13.330000	375.430000	NaN	NaN	1
75%	20000.000000	NaN	16.490000	567.300000	NaN	NaN	1
max	40000.000000	NaN	30.990000	1533.810000	NaN	NaN	1

11 rows × 27 columns


```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             396030 non-null float64
1   term                  396030 non-null object
2   int_rate              396030 non-null float64
3   installment           396030 non-null float64
4   grade                 396030 non-null object
5   sub_grade             396030 non-null object
6   emp_title             373103 non-null object
7   emp_length            377729 non-null object
8   home_ownership        396030 non-null object
9   annual_inc            396030 non-null float64
10  verification_status    396030 non-null object
11  issue_d               396030 non-null object
12  loan_status           396030 non-null object
13  purpose               396030 non-null object
14  title                 394275 non-null object
15  dti                   396030 non-null float64
16  earliest_cr_line      396030 non-null object
17  open_acc              396030 non-null float64
18  pub_rec               396030 non-null float64
19  revol_bal             396030 non-null float64
20  revol_util            395754 non-null float64
21  total_acc             396030 non-null float64
22  initial_list_status    396030 non-null object
23  application_type       396030 non-null object
24  mort_acc              358235 non-null float64
25  pub_rec_bankruptcies   395495 non-null float64
26  address               396030 non-null object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
```

✦ Correlation between independent features in the data set

```
plt.figure(figsize=(12,7))
sns.heatmap(df.corr(method='spearman'), annot=True, cmap='crest')
```

```
<ipython-input-10-64190c28e1eb>:2: FutureWarning: The default value of numeric_only i
sns.heatmap(df.corr(method='spearman'), annot=True, cmap='crest')
<Axes: >
```



Dropping "installement" features as it is related to "loan_amnt"


```
df.drop(columns=["installment"], inplace=True)
```

```
df.shape
```

```
(396030, 26)
```

```
plt.figure(figsize=(12,7))
sns.heatmap(df.corr(method='spearman'), annot=True, cmap="crest")
plt.show()
```

```
<ipython-input-13-50821c1d4bc5>:2: FutureWarning: The default value of numeric_only i
sns.heatmap(df.corr(method='spearman'), annot=True, cmap="crest")
```



```
df['loan_status'].value_counts(normalize=True) * 100

Fully Paid      80.387092
Charged Off     19.612908
Name: loan_status, dtype: float64
```

```
df.groupby(by='loan_status')['loan_amnt'].describe()
```

	count	mean	std	min	25%	50%	75%	n
loan_status								
Charged Off	77673.0	15126.300967	8505.090557	1000.0	8525.0	14000.0	20000.0	4000
Fully Paid	318357.0	13866.878771	8302.319699	500.0	7500.0	12000.0	19225.0	4000

There is a significant difference between the loan amount taken by charged_off and fully_paid users

```
df['home_ownership'].value_counts()

MORTGAGE      198348
RENT          159790
OWN           37746
OTHER         112
NONE          31
ANY           3
Name: home_ownership, dtype: int64
```

Most of the people applying for loans have either Mortgages or live in a rented house.

```
df.home_ownership.loc[(df.home_ownership == 'ANY') | (df.home_ownership == 'NONE')] = 'OTHER'
df.home_ownership.value_counts()
```

```
<ipython-input-17-50451bc6ac93>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus

```
df.home_ownership.loc[(df.home_ownership == 'ANY') | (df.home_ownership == 'NONE')] = 'OTHER'
MORTGAGE    198348
RENT        159790
OWN         37746
OTHER        146
Name: home_ownership, dtype: int64
```

Merged home_ownership of ANY, NONE and OTHER into a single group OTHER

```
# converting issue_d and earliest_cr_line to datetime data type
df["issue_d"] = pd.to_datetime(df["issue_d"])
df["earliest_cr_line"] = pd.to_datetime(df["earliest_cr_line"])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             396030 non-null float64
1   term                  396030 non-null object
2   int_rate              396030 non-null float64
3   grade                 396030 non-null object
4   sub_grade             396030 non-null object
5   emp_title              373103 non-null object
6   emp_length            377729 non-null object
7   home_ownership         396030 non-null object
8   annual_inc            396030 non-null float64
9   verification_status   396030 non-null object
10  issue_d               396030 non-null datetime64[ns]
11  loan_status           396030 non-null object
12  purpose               396030 non-null object
13  title                 394275 non-null object
14  dti                   396030 non-null float64
15  earliest_cr_line       396030 non-null datetime64[ns]
16  open_acc              396030 non-null float64
17  pub_rec               396030 non-null float64
18  revol_bal             396030 non-null float64
19  revol_util            395754 non-null float64
20  total_acc             396030 non-null float64
21  initial_list_status    396030 non-null object
22  application_type       396030 non-null object
23  mort_acc              358235 non-null float64
24  pub_rec_bankruptcies   395495 non-null float64
25  address               396030 non-null object
dtypes: datetime64[ns](2), float64(11), object(13)
memory usage: 78.6+ MB
```

```
df['title'] = df["title"].str.lower()
df["title"].value_counts()
```

```
debt consolidation    168108
credit card refinancing  51781
home improvement      17117
other                 12993
consolidation         5583
...
sweet                  1
mortgage conversion    1
debt consolidation and relocation  1
1 payment loan plan    1
toxic debt payoff      1
Name: title, Length: 41327, dtype: int64
```

Two important reason for taking loan is :

- debt consolidation
- credit card refinancing

```
def pub_rec(number):
    if number == 0.0:
        return 0
    else:
        return 1

def mort_acc(number):
    if number == 0.0:
        return 0
    else:
        return 1



def pub_rec_bankruptcies(number):
    if number == 0.0:
        return 0
    else:
        return 1

df['pub_rec'] = df.pub_rec.apply(pub_rec)
df['mort_acc'] = df.mort_acc.apply(mort_acc)
df['pub_rec_bankruptcies'] = df.pub_rec_bankruptcies.apply(pub_rec_bankruptcies)
df['loan_status'] = df.loan_status.map({'Fully Paid':0, 'Charged Off':1})
```

```
plt.figure(figsize=(17, 10))
plt.subplot(2,2,1)
grade= sorted(df["grade"].unique().tolist())
sns.countplot(data=df, x="grade", hue="loan_status", order=grade)

plt.subplot(2,2,2)
sub_grade= sorted(df["sub_grade"].unique().tolist())
sns.countplot(data=df, x="sub_grade", hue="loan_status", order=sub_grade)
```

<Axes: xlabel='sub_grade', ylabel='count'>



- Most loan belong to grade B and subgrade B3
- Loans with grade C are more likely to be charged off

```
plt.figure(figsize=(15,10))

plt.subplot(3,2,1)
sns.countplot(data=df, x='term', hue='loan_status')

plt.subplot(3,2,2)
sns.countplot(data=df, x='home_ownership', hue='loan_status')

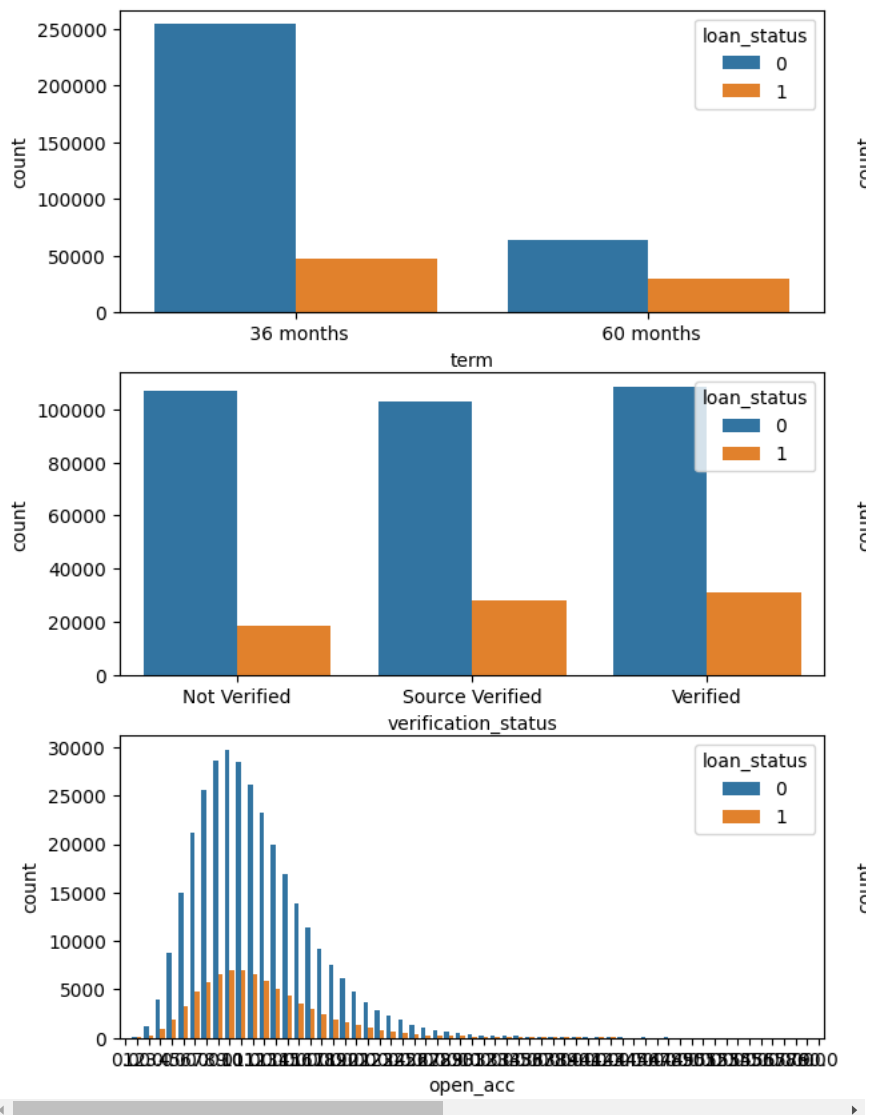
plt.subplot(3,2,3)
sns.countplot(data=df, x='verification_status', hue='loan_status')

plt.subplot(3,2,4)
sns.countplot(data=df, x='mort_acc', hue='loan_status')

plt.subplot(3,2,5)
sns.countplot(data=df, x='open_acc', hue='loan_status')

plt.subplot(3,2,6)
sns.countplot(data=df, x='pub_rec_bankruptcies', hue='loan_status')
```

<Axes: xlabel='pub_rec_bankruptcies', ylabel='count'>



```
df.emp_title.value_counts()[:3]
```

```
Teacher      4389
Manager      4250
Registered Nurse  1856
Name: emp_title, dtype: int64
```

Teachers and Managers are the most common applying for loan


```
df.emp_length.value_counts()
```

```
10+ years    126041
2 years      35827
< 1 year     31725
3 years      31665
5 years      26495
1 year       25882
4 years      23952
6 years      20841
7 years      20819
8 years      19168
9 years      15314
Name: emp_length, dtype: int64
```

```
plt.figure(figsize=(10,6))
```

```
order= [<" 1 year", "2 years", "3 years", "4 years", "5 years", "6 years", "7 years", "8 years", "9 years", "10+ years"]
sns.countplot(data=df, x='emp_length', hue="loan_status", order=order)
```

<Axes: xlabel='emp_length', ylabel='count'>




Higher the tenurity higher the chances of approving loan

```
bins= [0.0, 10000.0, 30000.0, 50000.0, 60000.0, 70000.0, 80000.0, 90000.0, 100000.0, 120000.0, 140000.0, 160000.0, 180000.0, 200000.0, 2:
label=["0_10k", "10k_30k", "30k_50k", "50k_60k", "60k_70k", "70k_80k", "80k_90k", "90k_1L", "1L_1.2L", "1.2L_1.4L", "1.4L_1.6L", "1.6L_1.8L", "1.8L_2L", "2L_2.25L", "2.25L_2.5L", "2.5L_2.75L", "2.75L_3L", "3L_3+"
df['binned_salary']= pd.cut(df["annual_inc"], bins=bins, labels=label)
```

```
plt.figure(figsize=(15, 6))
sns.countplot(data=df, x="binned_salary", hue="loan_status", order=label)
```

<Axes: xlabel='binned_salary', ylabel='count'>



Lesser the salary higher the chances of loan to be chared off and most loans applications are from people where salary range is between 10k to 1.2L

✓ Handling missing values

```
(df.isnull().sum()/len(df)) * 100
```



```

loan_amnt      0.000000
term           0.000000
int_rate       0.000000
grade          0.000000
sub_grade      0.000000
emp_title      5.789208
emp_length     4.621115
home_ownership 0.000000
annual_inc     0.000000
verification_status 0.000000
issue_d        0.000000
loan_status    0.000000
purpose        0.000000
title          0.443148
dti            0.000000
earliest_cr_line 0.000000
open_acc       0.000000
pub_rec        0.000000
revol_bal      0.000000
revol_util     0.069692
total_acc      0.000000
initial_list_status 0.000000
application_type 0.000000
mort_acc       0.000000
pub_rec_bankruptcies 0.000000
address        0.000000
binned_salary  0.000253
dtype: float64

```

```

total_acc_avg = df.groupby(by='total_acc').mean().mort_acc
def fill_mort_acc(total_acc, mort_acc):
    if np.isnan(mort_acc):
        return total_acc_avg[total_acc].round()
    else:
        return mort_acc
df['mort_acc'] = df.apply(lambda x: fill_mort_acc(x['total_acc'], x['mort_acc']), axis=1)

```

<ipython-input-30-5bbf87b41589>:1: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is deprecated. In a fut
total_acc_avg = df.groupby(by='total_acc').mean().mort_acc

```
(df.isnull().sum()/len(df)) * 100
```

```

loan_amnt      0.000000
term           0.000000
int_rate       0.000000
grade          0.000000
sub_grade      0.000000
emp_title      5.789208
emp_length     4.621115
home_ownership 0.000000
annual_inc     0.000000
verification_status 0.000000
issue_d        0.000000
loan_status    0.000000
purpose        0.000000
title          0.443148
dti            0.000000
earliest_cr_line 0.000000
open_acc       0.000000
pub_rec        0.000000
revol_bal      0.000000
revol_util     0.069692
total_acc      0.000000
initial_list_status 0.000000
application_type 0.000000
mort_acc       0.000000
pub_rec_bankruptcies 0.000000
address        0.000000
binned_salary  0.000253
dtype: float64

```

```

#droppig null values
df.dropna(inplace=True)
df.shape

```






```
(371126, 27)
```


✓ Handling outliers



```
num_data= df.select_dtypes(include='number')
num_col= num_data.columns
num_col

Index(['loan_amnt', 'int_rate', 'annual_inc', 'loan_status', 'dti', 'open_acc',
      'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc',
      'pub_rec_bankruptcies'],
      dtype='object')

for i in num_col:
    plt.figure(figsize=(6,4))
    sns.boxplot(df[i])
    plt.show()
```







```
for col in num_col:
    mean = df[col].mean()
    std = df[col].std()

    upper_limit = mean+3*std
    lower_limit = mean-3*std

    df = df[(df[col]<upper_limit) & (df[col]>lower_limit)]
```

```
df.shape

(355005, 27)
```

DATA PROCESSING

```
df['initial_list_status'].replace({'w': 0, 'f': 1 }, inplace=True)
```

```
df['initial_list_status'].value_counts()

1      214523
0      140482
Name: initial_list_status, dtype: int64
```

```
term_values = {' 36 months': 36, ' 60 months': 60}
df['term'] = df.term.map(term_values)
```

```
df["term"].value_counts()

36      270167
60      84838
Name: term, dtype: int64
```

```
df['zipcode'] = df.address.apply(lambda x: x[-5:])
```

```
df['zipcode']
```

```
0      22690
1      05113
2      05113
3      00813
4      11650
...
396025  30723
396026  05113
396027  70466
396028  29597
396029  48052
Name: zipcode, Length: 355005, dtype: object
```

```
#dropping columns
```

```
df.drop(columns=['issue_d', 'emp_title', 'title', 'sub_grade', 'address', 'earliest_cr_line', 'emp_length', ], inplace=True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 355005 entries, 0 to 396029
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  -
0   loan_amnt           355005 non-null  float64
1   term                355005 non-null  int64
2   int_rate            355005 non-null  float64
3   grade              355005 non-null  object
4   home_ownership      355005 non-null  object
5   annual_inc          355005 non-null  float64
6   verification_status 355005 non-null  object
7   loan_status         355005 non-null  int64
8   purpose            355005 non-null  object
9   dti                 355005 non-null  float64
10  open_acc            355005 non-null  float64
11  pub_rec             355005 non-null  int64
12  revol_bal           355005 non-null  float64
13  revol_util          355005 non-null  float64
14  total_acc           355005 non-null  float64
15  initial_list_status 355005 non-null  int64
16  application_type     355005 non-null  object
17  mort_acc            355005 non-null  int64
18  pub_rec_bankruptcies 355005 non-null  int64
19  binned_salary        355005 non-null  category
20  zipcode              355005 non-null  object
dtypes: category(1), float64(8), int64(6), object(6)
memory usage: 57.2+ MB
```

```
df.drop(columns=['binned_salary' ], inplace=True)
```

One Hot Encoding

```
dummies= ["grade", "home_ownership", "verification_status", "purpose", "application_type", "zipcode"]
df=pd.get_dummies(df, columns=dummies, drop_first=True)
```

```
df.shape
```

```
(355005, 49)
```

Buidling model

```
Y= df["loan_status"]
X= df.drop('loan_status', axis=1)
```

```
X_train, X_test, Y_train, Y_test= train_test_split(X, Y, test_size= 0.3, stratify=Y, random_state=42)
```

```
print(X_train.shape)
print(Y_train.shape)
```

```
(248503, 48)
(248503,)
```

```
#scaling
min_max_scaler= MinMaxScaler()
X_train= min_max_scaler.fit_transform(X_train)
X_test= min_max_scaler.transform(X_test)
```

```
lr=LogisticRegression(max_iter=1000)
lr.fit(X_train, Y_train)
```

```
▼ LogisticRegression
LogisticRegression(max_iter=1000)
```

```
from sklearn.metrics import accuracy_score, recall_score, precision_score, roc_auc_score, f1_score
```

```
print('Accuracy: ', accuracy_score(Y_test, lr.predict(X_test)))
```

```
Accuracy: 0.8907344463014779
```

```
Y_pred= lr.predict(X_test)
Y_pred
```

```
array([0, 0, 0, ..., 1, 1, 0])
```

```
confusion_matrixx= confusion_matrix(Y_test, Y_pred)
print(confusion_matrixx)
```


```
[[85448  562]
 [11075  9417]]
```

```
print(classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	86010
1	0.94	0.46	0.62	20492
accuracy			0.89	106502
macro avg	0.91	0.73	0.78	106502
weighted avg	0.90	0.89	0.88	106502

```
from sklearn.metrics import roc_curve, roc_auc_score
```

```
fpr, tpr, _ = roc_curve(Y_test, Y_pred)
auc = roc_auc_score(Y_test, Y_pred)
plt.plot(fpr, tpr, label="auc="+str(auc))
plt.plot([0, 1], [0, 1], 'r--')
plt.legend(loc=4)
plt.show()
```



Removing features which are multicollinear using Variance Inflation Factor

```
# VIF
from statsmodels.stats.outliers_influence import variance_inflation_factor
def calc_vif(X):
    # Calculating the VIF
    vif = pd.DataFrame()
    vif['Feature'] = X.columns
    vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by='VIF', ascending = False)
    return vif
```